

# Agenda

1. What & Why of HLD — High Level Design

support@scaler.com

2. Case Study - del.icio.us

3. Scaling challenges

4. HLD curriculum overview

Pragy → SSE + Instructor @ Scaler.

Media.Net - 2nd largest Ad Tech. - Platform Engineer.

Gatheringflow → web dev.  
30K+ new  
users each year

Predict what ads to show on what page.

1 Billion new row / day.

Past 20 days of data

~ 30 B rows of data (HBase)

MLE / Bandit Algos

{  
Machine Learning

Reinforcement Learning

UG → fin - 3 - MIET

PG → fin-1 - IITB

Programmers → write Code.  
↓ transition

Solving Problems.

/

Data Structures & Algo efficient.

Making sure that  
our problems stay  
solved

Low Level Design

Senior positions → Scaling Challenges.

Staff Engineer @ Google → 1<sup>+</sup> yr / annum

question during the design round.

Q: 'given a list of strings, sort them alphabetically'

[ Sneha ,  
Ganesh ,  
Saran ,  
Himanchu ,  
:  
]  
→ [ Ganesh ,  
Himanchu ,  
Saran ,  
Sneha ]

Python :

name = [...]  
sort(name)

Java: List <String> names;  
Collections.sort(names);

Q: "given a list of strings, sort them alphabetically"

Catch: There is 50 PB of data

1 bit  $\Rightarrow$  0/1

1 byte  $\Rightarrow$  8 bits

1 Kilo byte  $\Rightarrow$  1000 bytes =  $10^3$  bytes  
↳ SI prefix = 1000

1 KB

1 Kibbi Byte = 1024 bytes  
1 KiB  
Mibbi Byte / Gibbi Byte

1 Mega byte =  $10^6$  bytes

1 Giga byte =  $10^9$  bytes

50 PB = 50,000,000 GB

1 Tera byte =  $10^{12}$  bytes

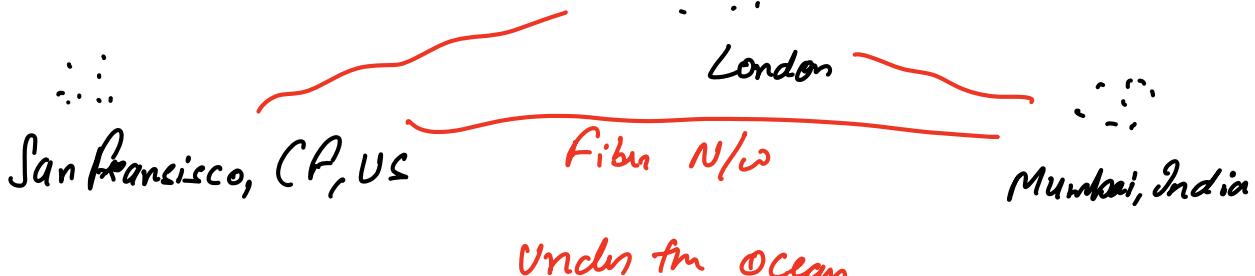
1 Peta byte =  $10^{15}$  bytes

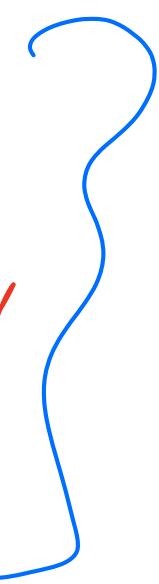
50 PB fits in 1 system? No!

avg laptop has 1 TiB Harddisk.

you need 50,000 laptops to just store this data.

Massive Distributed System.



- Despit fn issue.
- ① N/w can go down
  - ② Machines can fail
  - ③ Data can get corrupted
  - ④ Natural disasters
  - ⑤ Network latency
- amount of time taken for something to happen.  
"delay"
- 

## Map - Reduce

Simple problem → Scale  
v. v. difficult

What challenges appear at scale, & how to tackle them.

## Pre-requisites

- ① Writing Code
- ② Data Structure & Algo
- ③ Computer Fundamentals
  - └ Networks
  - └ Databases (SQL)
  - └ Operating Systems.
- ④ Low Level Design

Mastered

90% understanding of the concepts discussed today → ok ✓



# Del.icio.us Case Study

Joshua

↳ bookmarking Service in the cloud.

? Launched in



2002



2005



2006



2008

(2Kb/s speed)

Dial-Up

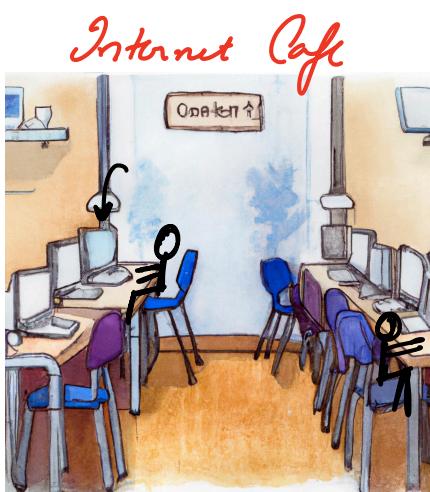
↓

Broad Band

↓

Fiber 1 Gbps

## Motivation



Internet Cafe

google.com / scalen.com /  
---

Bookmark them

→ no bookmarks  
in other PC

## Features

Minimal Viable Product (MVP) - Proof of Concept

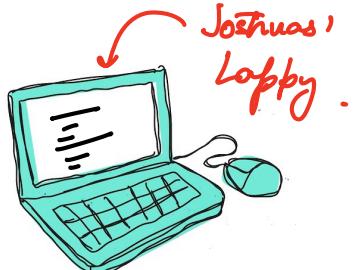
1. Auth - User registration & login

2. Add Bookmark

add Bookmark (user-id, URL)

3. Fetch Bookmarks

get My Bookmarks (user-id)



Python + Django + MySQL  
Web App

127.0.0.1:8000/

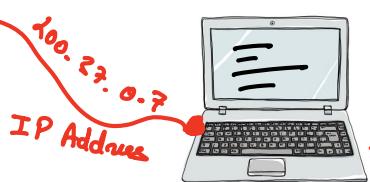
to be continued..



# Fantastic websites & where to find them

ISP → Internet Service Provider

Joshua's Laptop



100.37.0.7  
IP Address

Ashok's Laptop



find me the service running  
at 100.37.0.7 → Joshua's IP.



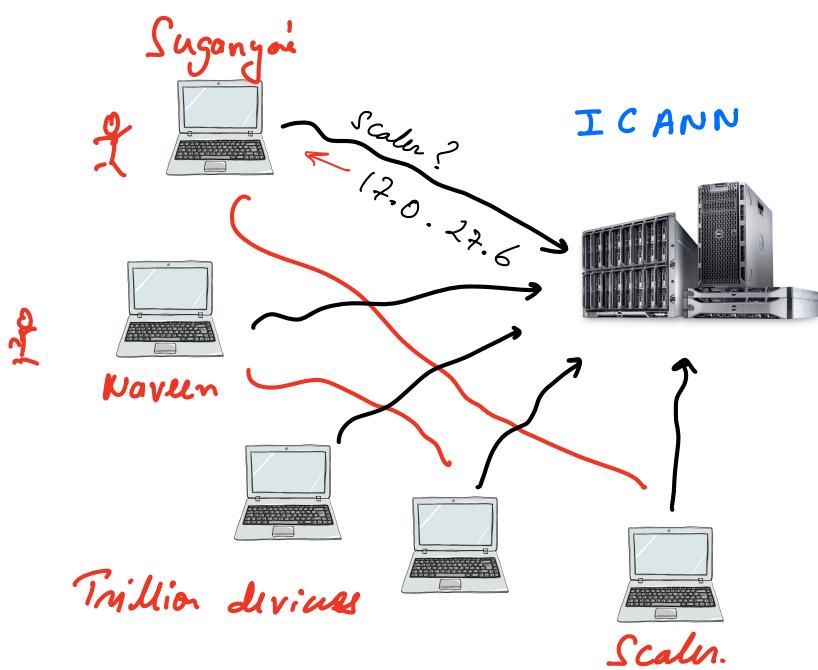
Quiz - how does the browser know?

DNS Service

<http://delicious.us>

Domain Name Mapping

Domain Name	IP Address
delicious.com	17.0.27.6
google.com	200.70.20.10
delicious	12.10.1.2
--	



ICANN → I CAN'T!!

Singl Point of failure  
SPOF

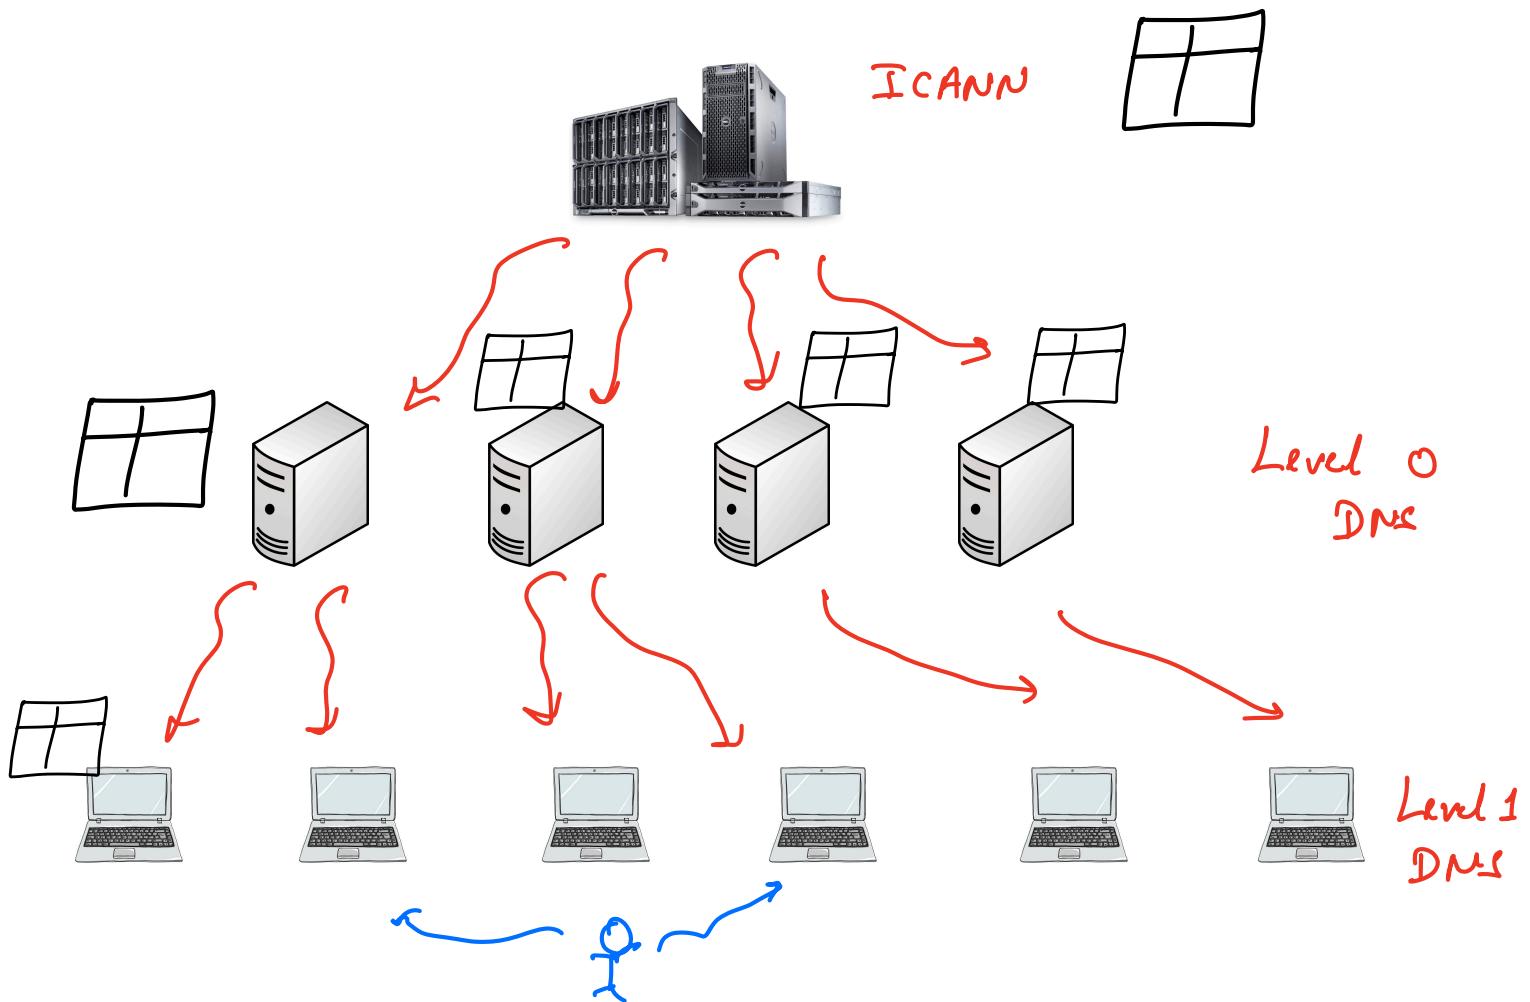


Bottle neck.



# Domain Name Service

<https://www.cloudflare.com/en-gb/learning/dns/what-is-dns/>



Who maintains the DNS?

Internet Companies → Google / Microsoft / --

IISPs → Airtel / Jio / Comcast

Educational Institutes → Stanford / Harvard  
Military

Anyone can  
create a DNS  
server.

Demos



# Del.icio.us

continued..

2006 → 25,000 Rs

Intel Pentium Core 2 duo → 2 cores

128 MB of RAM

40 GB of HDD

@ 2.2 GHz

*Joshua's System*



RAM

128 MB

Disk

40 GB

CPU

2 core

25,000 Rs

System Configuration



*App      DB      OS*



DB



OS



Django + Python      MySQL      Apache



Issues

→ if Joshua is running the website on his personal machine.

① Down time — sleep / powercut / internet issues

② Content for user can b/w website & personal usage

Dedicated Machine to serve the website → **Server**.

↓  
Machine connected to internet  
with the sole purpose of providing  
some service



Quiz - space

usn-bookmarks

usn-id	URL
1	google.com
2	Scalari.com
1	Youtube.com
2	microsoft.com

8 bytes

1000 bytes

Size of each row is 1 KB

1 million new bookmarks each day.

1 bit  $\Rightarrow$  0 or 1

fundamental unit of information.

1 byte  $\Rightarrow$  8 bits

1 KB  $\Rightarrow$   $10^3$  bytes

1 MB  $\Rightarrow$   $10^6$  bytes

1 GB  $\Rightarrow$   $10^9$  bytes

Every bookmark  $\rightarrow$  row in the database

Size of 1 row = 1 KB =  $10^3$  bytes / bookmark

New Bookmarks each day = 1 Million  
=  $10^6$  bookmarks / day.

How much data are we getting / day?

=  $10^6$  bookmarks / day. \*  $10^3$  bytes / bookmark

=  $10^9$  bytes / day. = 1 GB / day

Start with 40 GB space

Occupying 1 GB / day.

# of days for which we can survive =  $\frac{40 \text{ GB}}{1 \text{ GB/day}} = 40 \text{ days}$



Quiz - time

$$10^2 \text{ Bytes/gb} * 10^6 \text{ gb/day}$$

1 GB / day.

$$= 10^9 \text{ bytes/day} = 1 \text{ GB/day.}$$

90 GB

$$\# \text{ days} = \frac{90 \text{ GB}}{1 \text{ GB/day}} = 90 \text{ days}$$

10:43 → 10:55



## Solutions

### Buy Powerful Machines

1.



15,000 Rs

128 MB RAM

40 GB HDD

2 core CPU

}

40 days

moved

Local vendor → give me the best system  
→ cash! your hand!



2 Lakh Rs

512 MB RAM

100 GB HDD

2 core CPU

→ 60 additional GB

60 more days



IBM → Seven grade hardware

20 Lakh Rs

2 GB RAM

500 GB HDD

8 core CPU

→ 400 additional days.



Seven Rack

2 Core Rs

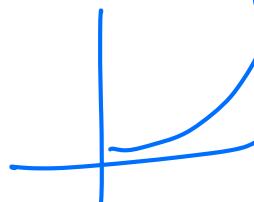
× 1 year

10 GB RAM

→ 5 more years

2 TB HDD

22 core CPU



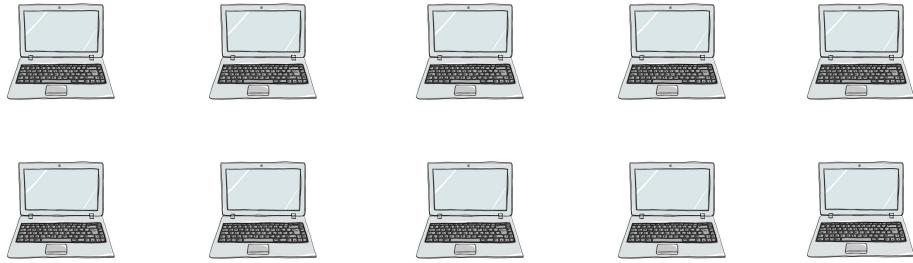
## 2. Buy lots of cheap machines

15000 Rs

128MB RAM

40GB HDD

2 core CPU



2 core Rs  $\Rightarrow$  570

effective RAM: 72GB

HDD: 22TB

Cores: 1000<sup>+</sup>

80% discount!

Economy of scale

Vertical

Horizontal

① Super Easy  $\rightarrow$  throw money

① v.v. difficult  
entire HLD is about solving these difficulties.

② Hard Limit  
based on current day technology

② Do this infinitely !!

③ New machine  $\rightarrow$  stop using old machine

③ Buy new machine & continue using old one.

Eventually you will need Horizontal Scaling.

Potential of vertical Scaling nowadays.

## Typical Server Config

RAM: 2 GB → 128 GB

HDD: 512 GiB → 8 TB

CPU : 1 core → 8 cores

N/w : 10Mbps → 1 Gbps

Max server config  
as of Sept 2023

RAM: 12 TB

HDD: 2 PB

CPU : 570 cores  
1GB of L2 cache

570

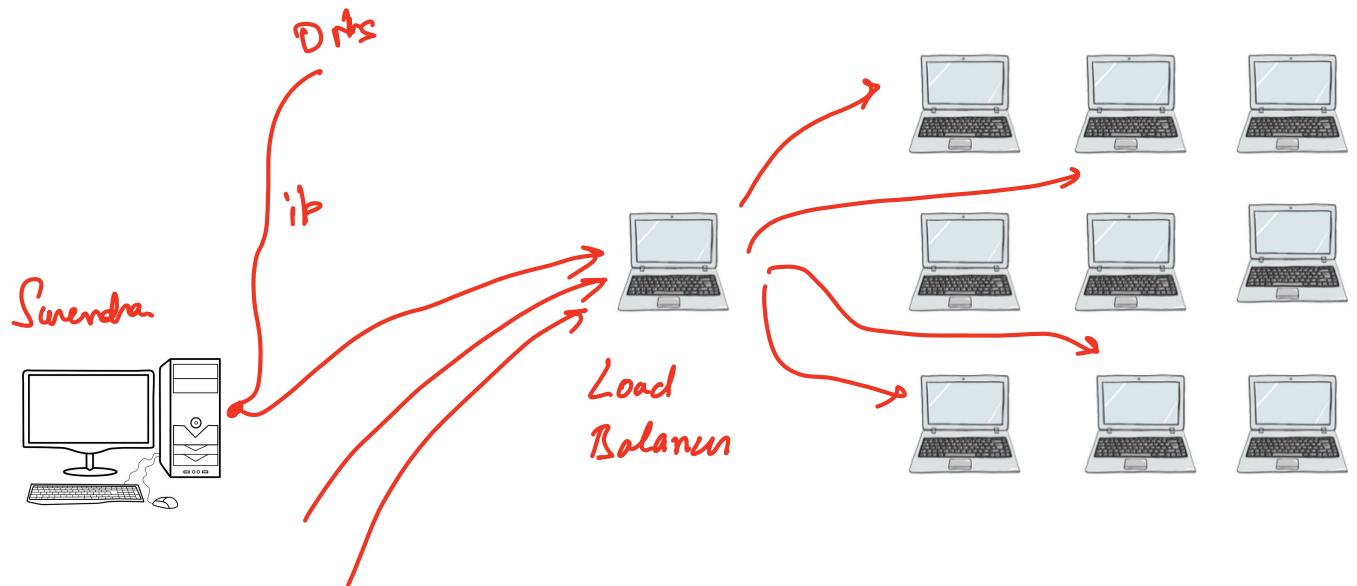
cores

EPYC

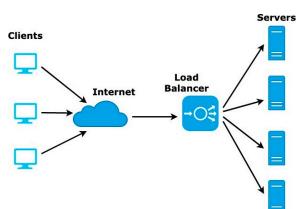


Quiz - which ip?

500



## Home work



## Load Balancing

[https://docs.google.com/document/d/1DxQzLpu1XPe\\_mRWsewNWtKL6E4uwKHQhBp7GX6Sg7qI/edit](https://docs.google.com/document/d/1DxQzLpu1XPe_mRWsewNWtKL6E4uwKHQhBp7GX6Sg7qI/edit)



How does the Load Balancer track which servers are up & running?





# Further Challenges

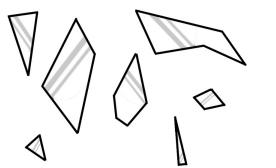
?

What if the Load Balancer goes down?

?

Which machine should we send the request to?  
Routing Algorithm

<https://docs.nginx.com/nginx/admin-guide/load-balancer/http-load-balancer/>



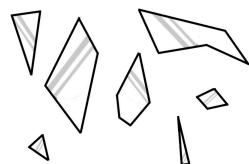
How to store the data?

?

Can we store all data on 1 machine?

?

Split randomly?



Sharding



*Gangnam Style* → cross billion views  
 crossed 2 billion views.

view count → (-ve) → 4 billion diff values

32 bit signed integer →  $[-2^{31} \dots 2^{31}-1]$

$\approx [-2\text{ Billion} \dots 2\text{ Billion}]$

World Pop =

$$2^3 \cdot 2^3 \cdot 2 = 8 \text{ choices}$$

0/1   0/1   0/1        3 bits

32 bits

$2^{32}$  choices ≈ 4 billion

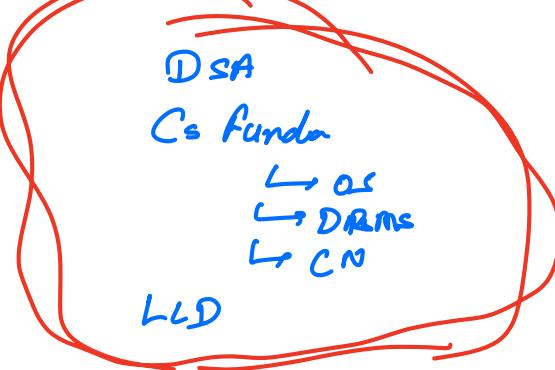
0 0 0  
0 0 1  
0 1 0  
0 1 1  
1 0 0  
1 0 1  
1 1 0  
1 1 1

8 =  $2^3$

8 billion prob

# HL D curriculum

1.5 month



- ① Load Balancing & Consistency Hashing.
- ② Caching → Local / global / distributed  
→ Eviction  
→ Invalidity → write thru / write-back / ETZ

Case studies → Fb news feed cached

Codjudge @ stan

Undercover @ columbia

- ③ CAP theorem / PACELC Theorem / Eventual Consistency
- ④ Master - slave replication → Quorum  
→ Turntable consistency  
→ Lead Election
- ⑤ SQL vs NoSQL db.

ACID / BASE

Key-value      Redis / DynamoDB

Doc            MongoDB / -

Column        Cassandra / Hbase

How to choose

⑥ Sharding → how to shard

⑦ How NoSQL db work internally. → LSM Trees.  
 Bloom Filter

⑧ Case Study → Design Google Typewriter.

⑨ Case Study → Fb Messenger.

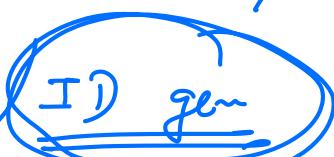
⑩ Case Study → Zookeeper

⑪ Messaging Queue → Kafka

⑫ Case Study — Elastic Search

⑬ Case Study — S3 / file storage

⑭ Case Study — Ubs / nearest neighbor → Quad Tree

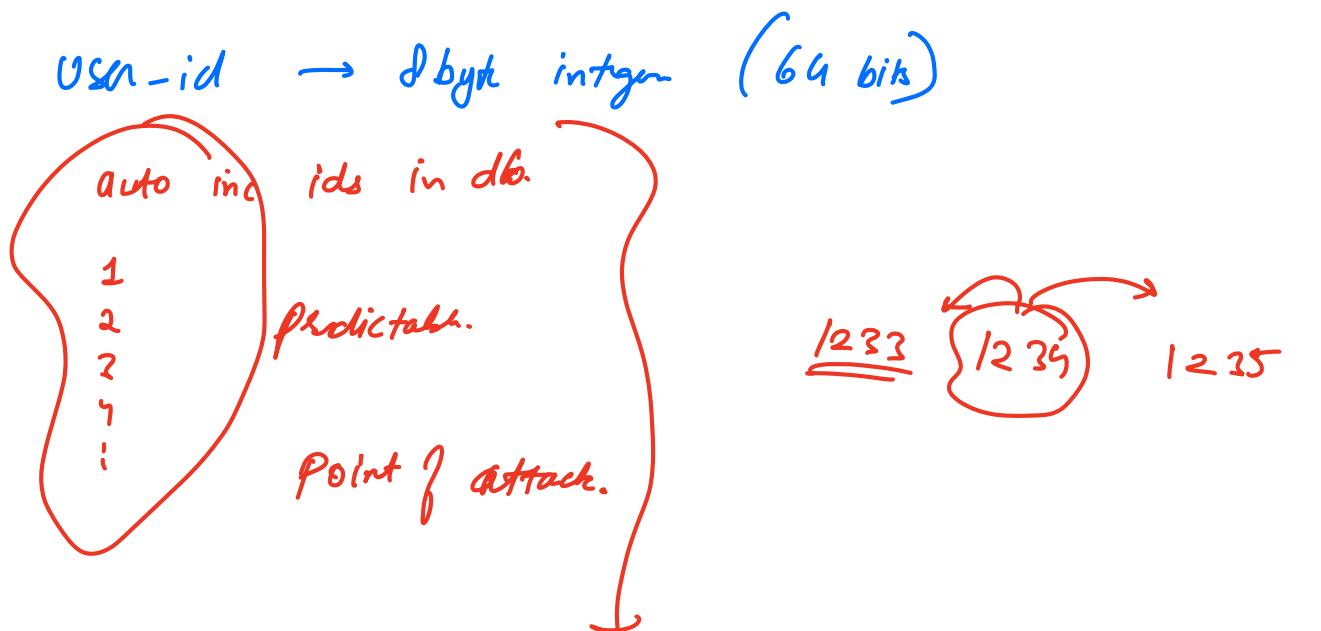
⑮ Rate Limiter 

⑯ Hot spot

⑰ Microservices — Monolithic / Event Driven / Pub Sub /  
Saga / Dist. Boxes

⑱ Case study → IRCTC / ticket booking.

Stack overflow  $\rightarrow$  SQL db  
100 TR of in-vm (RAM) cache

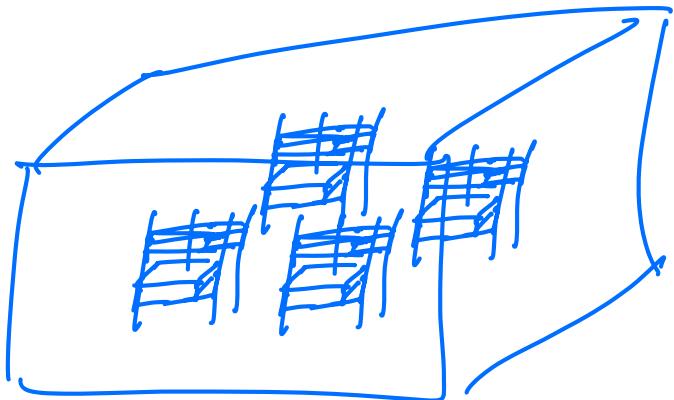


UUID / GUID  
Random.  
a72b6 - 3hab3 - aff 26255 - a --  
16 byte value  
16 bytes = 128 bits no.

$$(23)_{10} = (10111)_2 = (17)_{16} = (27)_8$$

Shoe Rack stores shoes

Shows rack stores shows.



Cloud - someone else's computer. over the n/w.